

## 2.3 Rolling Validation for Multiple Horizons

Due to the dependencies of instances near each other in time,  $k$ -fold cross-validation will not work for time-series data. A simple form of rolling validation divides a dataset into an initial training set and test set. For example, we can use the first 60% of the samples as the training set and rest as the test set. For horizon  $h = 1$  forecasting, the first value in the test set is forecasted based on the model produced by training on the training set. The error is the difference between the actual value in the test set and the forecasted value which is used to calculate the symmetric mean absolute percentage error (sMAPE). sMAPE is used as the primary performance metric, which is one of the standard performance metrics in time-series forecasting. Although, other performance metrics such as MAPE can also be used.

$$\text{sMAPE} = 100 * \text{mean}(2|y_t - f_t|/|y_t| + |f_t|)$$

where  $y_t$  is the true valued time-series vector and  $f_t$  is the forecasted time-series vector.

For simple, efficient models, the process of rolling forward to forecast the next value in the test set often involves retraining the model by including the first value from the test set into the training set (at the end) and discarding the first value from the training set. To maintain the same size for the training set, we remove the first value from the training set. We adjust the frequency of retraining for the statistical models such that we forecast  $kt$  samples ahead in the test set before including them in the training set and retraining our model.

### 2.3.1 Rolling Validation in R

The rolling validation function in R is used to generate multi-horizon out-of-sample forecasts for a given time-series. The *forecast* package developed by Dr. Rob J. Hyndman is used to get the ARIMA model in R. The entire code is given at the end and we begin with the stepwise explanation for the code below:

We begin by first installing and using the *forecast* package in our program. Package Installation is required only for first time setup and can be commented out for later runs.

```
1 install.packages("forecast")
2 library(forecast)
```

Listing 2.11: Step 1

We then begin with our initial variable setup for our program. This includes training/test size, horizons, all the hyperparameters of our ARIMA/SARIMA model and the rolling validation function. The variable `zero_reference` is used as first day starts on 1<sup>st</sup> Jan 1970 in R language and until 25<sup>th</sup> Feb 2020, which is the start of our model data, the number of days is equal to 18317.

```
1 zero_reference = 18317 #Difference between 1970-01-01 to 2020-02-25
2
3 ts_size = 376 # Number of Sample/Timesteps in entire time series
4 tr_size = 222 # Number of Samples/Timesteps in Training set
```

```

5 te_size = 154 # Number of Samples/Timesteps in Training set
6 horizons = 14 # Number of days to forecast ahead in the future
7 kt = 1 # Retraining Frequency (No. of Samples)
8 rt = 0 # Total Number of Retrainings Occured
9 p = 1 # Non-seasonal Auto-Regressive Order
10 d = 0 # Non-seasonal Differencing
11 q = 0 # Non-seasonal Moving Average Order
12 PP = 3 # Seasonal Auto-Regressive Order
13 DD = 1 # Seasonal Differencing
14 QQ = 1 # Seasonal Moving Average Order
15 s = 7 # Seasonal Period

```

Listing 2.12: Step 2

Once we have our initial setup ready, we then pull our dataset from remote URL, in this case GitHub repo, and then get the response variable, in this case deathIncrease, and finally convert it into a zoo object. A zoo object is used to store and easily manipulate regular and irregular time-series vectors.

```

1 # Pull in the dataset from given URL
2 covid_data <- read.csv(file = "https://raw.githubusercontent.com/scalation
  /data/master/COVID/CLEANED_35_Updated.csv")
3 #Convert reponse column into "zoo" object
4 response_variable = covid_data$deathIncrease
5 #@see: stackoverflow.com/questions/33714660/what-is-the-difference-the-zoo
  -object-and-ts-object-in-r
6 covid_ts <- zoo(response_variable, seq(as.Date("2020-01-13"), as.Date("
  2021-03-07"), by = "days"), frequency = 1)

```

Listing 2.13: Step 3

Once we have our dataset ready, we then build our forecast matrix for out-of-sample forecast. Denoting forecast matrix as  $f_m$ , we have  $f_m \in R^{h \times te\_size}$ , where  $h$  is the number of horizons we are forecasting for and  $te\_size$  is the size of the test set.

```

1 #Build forecasting matrix to hold all rolling forecasts. Matrix Dimension:
  horizons by size of test set
2 forecast_matrix <- matrix(data = c(-1.2), nrow = horizons, ncol = te_size)

```

Listing 2.14: Step 4

We are now ready to begin with the rolling validation, which is done by looping over the test set as given below:

```

1 # Start rolling validation: run through entire test set 1 sample at a time
2 for (i in 1:te_size) {
3
4   # Get training data: start=1+18317 from 1970-01-01 which is 2020-02-26
   & end=221+18317+1 which is 2020-10-04(inclusive)
5   # This also shifts the training window ahead by 1 sample at a time
   into the test set on every every iteration

```

```

6   # For eg. For the next iteration: start=2+18317 from 1970-01-01 which
   is 2020-02-27 & end=221+18317+2 which is 2020-10-05(inclusive)
7   # So we dropped the sample for 2020-02-26 from the training in 2nd
   iteration and included 1 from the test set. That's how we roll forward.
8   model_data <- window(covid_ts, start = as.Date(i + zero_reference),
   end = as.Date(221 + zero_reference + i), frequency = 1)
9
10  # Retraining after kt samples
11  if (i == 1) {
12      arima_model <- Arima(y = model_data, order = c(p, d, q), seasonal
   = list(order = c(PP, DD, QQ), period = s), method = "ML")
13  } else {
14      if ((i %% kt) == 0) { #if iteration mod retrain frequency is 0,
   then it is time to retrain the model with new training data
15          rt = rt + 1
16          arima_model <- Arima(y = model_data, order = c(p, d, q),
   seasonal = list(order = c(PP, DD, QQ), period = s), method = "ML")
17      } else {
18          # if not, then
19          # 1) use the new shifted data as the model data and
20          # 2) do not reestimate the model parameters again: done by
   specifying the old model reference variable in "model" attribute of the
   function call
21          arima_model <- Arima(y = model_data, model = arima_model)
22      }
23  }
24  # Once we have the model set: either with new data or with new data
   and retrained get next "horizon" number of forecasts.
25  future <- forecast(arima_model, h = horizons, level = 0)
26  # Get the single point forecast vector as forecast() function also
   returns high/max and low/min values
27  forecasts <- future$mean
28  # Start filling the forecast matrix 1 by 1 for "horizon" number of
   forecasts
29  for (k in 1:horizons) {
30      if (k + i - 1 <= te_size) { # Condition to prevent overflow at the
   end of the test set;
31          forecast_matrix[k, k+i-1] = forecasts[k]
32      }
33  }
34 }

```

Listing 2.15: Step 5

After the rolling validation is complete, we have our out-of-sample forecasts collected in the forecasting matrix but the initial  $h$  number of days for the corresponding  $h^{th}$  horizon does not have any forecast because we cannot generate  $h = 2$  forecast for day 1 or  $h = 3$  forecast for day 2 and so on. In order to tackle this issue, we use the previous horizon forecast by copying the latest generated value into all the

remaining rows for that single day. Once that is done we store the transpose of the matrix where we get each horizon forecast as column vectors.

```

1 # For first 14 days copy the values of the previous day as we cannot have
  h=2 forecast for first day of test set, h=3 forecast for the 2nd day of
  test set and so on;
2 for(i in 2:horizons){
3   for(j in 1:i){
4     forecast_matrix[i,j] = forecast_matrix[i-1,j]
5   }
6 }
7 # Transpose the forecast matrix to be of dimension: te_size by horizons
8 forecast_matrix <- t(forecast_matrix)

```

Listing 2.16: Step 6

We now get our true values for the test set

```

1 # Get true values of the test data
2 test_data <- window(covid_ts, start = as.Date((tr_size+1) + zero_reference
  ), end = as.Date(ts_size + zero_reference), frequency = 1)

```

Listing 2.17: Step 7

Next, we build the matrix to store SMAPE scores for all horizons.

```

1 # Generate matrix to store smape scores for each horizon
2 smape_scores = matrix(data = NA, nrow = horizons, ncol = 1)

```

Listing 2.18: Step 8

Finally, we loop through each column in our forecasting matrix and compare it with true values of the test set to calculate smape score and store it in our smape score matrix.

```

1 # Calculate the smape scores and store in the smape matrix, horizon by
  horizon.
2 for (i in 1:horizons) {
3   smape_scores[i, 1] <- mean(abs((test_data - forecast_matrix[, i]) / (
    test_data + forecast_matrix[, i]))) * 200 #smape formula
4   print(paste0("h=",i," ,smape=",smape_scores[i,1]))
5 }

```

Listing 2.19: Step 9

Below is the the full code for all above steps combined. The code is also available on github at: [https://github.com/scalation/data/blob/master/COVID/covid\\_forecasting.r](https://github.com/scalation/data/blob/master/COVID/covid_forecasting.r)

```

1 #install.packages("forecast") #Unccomment and run this for the first time
  R is setup on your machine
2 library(forecast)
3
4 zero_reference = 18317 #Difference between 1970-01-01 to 2020-02-25
5

```

```

6 ts_size = 376 # Number of Sample/Timesteps in entire time series
7 tr_size = 222 # Number of Samples/Timesteps in Training set
8 te_size = 154 # Number of Samples/Timesteps in Training set
9 horizons = 14 # Number of days to forecast ahead in the future
10 kt = 1 # Retraining Frequency (No. of Samples)
11 rt = 0 # Total Number of Retrainings Occured
12 p = 1 # Non-seasonal Auto-Regressive Order
13 d = 0 # Non-seasonal Differencing
14 q = 0 # Non-seasonal Moving Average Order
15 PP = 3 # Seasonal Auto-Regressive Order
16 DD = 1 # Seasonal Differencing
17 QQ = 1 # Seasonal Moving Average Order
18 s = 7 # Seasonal Period
19
20 # Pull in the dataset from given URL
21 covid_data <- read.csv(file = "https://raw.githubusercontent.com/scalation
/data/master/COVID/CLEANED_35_Updated.csv")
22
23 #Convert reponse column into "zoo" object
24 response_variable = covid_data$deathIncrease
25 #@see: stackoverflow.com/questions/33714660/what-is-the-difference-the-zoo
-object-and-ts-object-in-r
26 covid_ts <- zoo(response_variable, seq(as.Date("2020-01-13"), as.Date("
2021-03-07"), by = "days"), frequency = 1)
27
28 #Build forecasting matrix to hold all rolling forecasts. Matrix Dimension:
horizons by size of test set
29 forecast_matrix <- matrix(data = c(-1.2), nrow = horizons, ncol = te_size)
30
31 # Start rolling validation: run through entire test set 1 sample at a time
32 for (i in 1:te_size) {
33
34 # Get training data: start=1+18317 from 1970-01-01 which is 2020-02-26
& end=221+18317+1 which is 2020-10-04(inclusive)
35 # This also shifts the training window ahead by 1 sample at a time
into the test set on every every iteration
36 # For eg. For the next iteration: start=2+18317 from 1970-01-01 which
is 2020-02-27 & end=221+18317+2 which is 2020-10-05(inclusive)
37 # So we dropped the sample for 2020-02-26 from the training in 2nd
iteration and included 1 from the test set. That's how we roll forward.
38 model_data <- window(covid_ts, start = as.Date(i + zero_reference),
end = as.Date(221 + zero_reference + i), frequency = 1)
39
40 # Retraining after kt samples
41 if (i == 1) {
42 arima_model <- Arima(y = model_data, order = c(p, d, q), seasonal
= list(order = c(PP, DD, QQ), period = s), method = "ML")

```

```

43   } else {
44       if ((i %% kt) == 0) { #if iteration mod retrain frequency is 0,
then it is time to retrain the model with new training data
45           rt = rt + 1
46           arima_model <- Arima(y = model_data, order = c(p, d, q),
seasonal = list(order = c(PP, DD, QQ), period = s), method = "ML")
47       } else {
48           # if not, then
49           # 1) use the new shifted data as the model data and
50           # 2) do not reestimate the model parameters again: done by
specifying the old model reference variable in "model" attribute of the
function call
51           arima_model <- Arima(y = model_data, model = arima_model)
52       }
53   }
54   # Once we have the model set: either with new data or with new data
and retrained get next "horizon" number of forecasts.
55   future <- forecast(arima_model, h = horizons, level = 0)
56   # Get the single point forecast vector as forecast() function also
returns high/max and low/min values
57   forecasts <- future$mean
58   # Start filling the forecast matrix 1 by 1 for "horizon" number of
forecasts
59   for (k in 1:horizons) {
60       if (k + i - 1 <= te_size) { # Condition to prevent overflow at the
end of the test set;
61           forecast_matrix[k, k+i-1] = forecasts[k]
62       }
63   }
64 }
65 # For first 14 days copy the values of the previous day as we cannot have
h=2 forecast for first day of test set, h=3 forecast for the 2nd day of
test set and so on;
66 for(i in 2:horizons){
67     for(j in 1:i){
68         forecast_matrix[i,j] = forecast_matrix[i-1,j]
69     }
70 }
71 # Transpose the forecast matrix to be of dimension: te_size by horizons
72 forecast_matrix <- t(forecast_matrix)
73
74 # Get true values of the test data
75 test_data <- window(covid_ts, start = as.Date((tr_size+1) + zero_reference
), end = as.Date(ts_size + zero_reference), frequency = 1)
76
77 # Generate matrix to store smape scores for each horizon
78 smape_scores = matrix(data = NA, nrow = horizons, ncol = 1)

```

```

79
80 # Calculate the smape scores and store in the smape matrix, horizon by
    horizon.
81 for (i in 1:horizons) {
82   smape_scores[i, 1] <- mean(abs((test_data - forecast_matrix[, i]) / (
      test_data + forecast_matrix[, i]))) * 200 #smape formula
83   print(paste0("h=", i, ", smape=", smape_scores[i, 1]))
84 }

```

Listing 2.20: Full Code